

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

from sklearn.preprocessing import StandardScaler
from sklearn.neural_network import MLPRegressor
from sklearn.neural_network import MLPClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_absolute_error
```

```
In [2]: #Fetch data from file
my_data = open(r"C:\\Users\\user\\python_example\\HomeC.csv")
df = pd.read_csv(my_data)
df
```

C:\Users\user\AppData\Local\Temp\ipykernel_9040\188017517.py:3: DtypeWarning: Columns (0,27) have mixed types. Specify dtype option on import or set low_memory=False.

```
df = pd.read_csv(my_data)
```

Out[2]:

	time	use [kW]	gen [kW]	House overall [kW]	Dishwasher [kW]	Furnace 1 [kW]	Furnace 2 [kW]	Home office [kW]	
0	1451624400	0.932833	0.003483	0.932833	0.000033	0.020700	0.061917	0.442633	0.
1	1451624401	0.934333	0.003467	0.934333	0.000000	0.020717	0.063817	0.444067	0.
2	1451624402	0.931817	0.003467	0.931817	0.000017	0.020700	0.062317	0.446067	0.
3	1451624403	1.022050	0.003483	1.022050	0.000017	0.106900	0.068517	0.446583	0.
4	1451624404	1.139400	0.003467	1.139400	0.000133	0.236933	0.063983	0.446533	0.
...
503906	1452128306	1.599333	0.003233	1.599333	0.000050	0.104017	0.625033	0.041750	0.
503907	1452128307	1.924267	0.003217	1.924267	0.000033	0.422383	0.637733	0.042033	0.
503908	1452128308	1.978200	0.003217	1.978200	0.000050	0.495667	0.620367	0.042100	0.
503909	1452128309	1.990950	0.003233	1.990950	0.000050	0.494700	0.634133	0.042100	0.
503910	\	NaN	NaN	NaN	NaN	NaN	NaN	NaN	

503911 rows × 32 columns



In [3]: `df.head(10)`

Out[3]:

	time	use [kW]	gen [kW]	House overall [kW]	Dishwasher [kW]	Furnace 1 [kW]	Furnace 2 [kW]	Home office [kW]	Fridg [kW]
0	1451624400	0.932833	0.003483	0.932833	0.000033	0.020700	0.061917	0.442633	0.12415
1	1451624401	0.934333	0.003467	0.934333	0.000000	0.020717	0.063817	0.444067	0.12400
2	1451624402	0.931817	0.003467	0.931817	0.000017	0.020700	0.062317	0.446067	0.12353
3	1451624403	1.022050	0.003483	1.022050	0.000017	0.106900	0.068517	0.446583	0.12313
4	1451624404	1.139400	0.003467	1.139400	0.000133	0.236933	0.063983	0.446533	0.12285
5	1451624405	1.391867	0.003433	1.391867	0.000283	0.503250	0.063667	0.447033	0.12230
6	1451624406	1.366217	0.003450	1.366217	0.000283	0.499400	0.063717	0.443267	0.12205
7	1451624407	1.431900	0.003417	1.431900	0.000250	0.477867	0.178633	0.444283	0.12180
8	1451624408	1.627300	0.003417	1.627300	0.000183	0.447650	0.365700	0.441467	0.12161
9	1451624409	1.735383	0.003417	1.735383	0.000017	0.171550	0.682500	0.438733	0.12163

10 rows × 32 columns

In [4]: `df.tail(10)`

503904	1452128304	1.608867	0.003217	1.608867	0.000033	0.114300	0.623283	0.041817	
503905	1452128305	1.601233	0.003183	1.601233	0.000050	0.085267	0.642417	0.041783	
503906	1452128306	1.599333	0.003233	1.599333	0.000050	0.104017	0.625033	0.041750	
503907	1452128307	1.924267	0.003217	1.924267	0.000033	0.422383	0.637733	0.042033	
503908	1452128308	1.978200	0.003217	1.978200	0.000050	0.495667	0.620367	0.042100	
503909	1452128309	1.990950	0.003233	1.990950	0.000050	0.494700	0.634133	0.042100	
503910	\	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN

10 rows × 32 columns

```
In [5]: # this will describe all the numeric values on the dataset
df.describe()
```

Out[5]:

	use [kW]	gen [kW]	House overall [kW]	Dishwasher [kW]	Furnace 1 [kW]	Furnace 2 [kW]
count	503910.000000	503910.000000	503910.000000	503910.000000	503910.000000	503910.000000
mean	0.858962	0.076229	0.858962	0.031368	0.099210	0.136710
std	1.058207	0.128428	1.058207	0.190951	0.169059	0.178610
min	0.000000	0.000000	0.000000	0.000000	0.000017	0.000000
25%	0.367667	0.003367	0.367667	0.000000	0.020233	0.064410
50%	0.562333	0.004283	0.562333	0.000017	0.020617	0.066610
75%	0.970250	0.083917	0.970250	0.000233	0.068733	0.080610
max	14.714567	0.613883	14.714567	1.401767	1.934083	0.794910

8 rows × 28 columns



In [6]: *#Check for missing values on the dataset.if zero corresponds with that particu*
df.isnull().sum()

Out[6]:

time	0
use [kW]	1
gen [kW]	1
House overall [kW]	1
Dishwasher [kW]	1
Furnace 1 [kW]	1
Furnace 2 [kW]	1
Home office [kW]	1
Fridge [kW]	1
Wine cellar [kW]	1
Garage door [kW]	1
Kitchen 12 [kW]	1
Kitchen 14 [kW]	1
Kitchen 38 [kW]	1
Barn [kW]	1
Well [kW]	1
Microwave [kW]	1
Living room [kW]	1
Solar [kW]	1
temperature	1
icon	1
humidity	1
visibility	1
summary	1
apparentTemperature	1
pressure	1
windSpeed	1
cloudCover	1
windBearing	1
precipIntensity	1
dewPoint	1
precipProbability	1
dtype: int64	

In [7]: *#This shows you the data types of the values assigned to particular variable*
df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 503911 entries, 0 to 503910
Data columns (total 32 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   time                                503911 non-null  object
1   use [kW]                            503910 non-null  float64
2   gen [kW]                            503910 non-null  float64
3   House overall [kW]                  503910 non-null  float64
4   Dishwasher [kW]                    503910 non-null  float64
5   Furnace 1 [kW]                      503910 non-null  float64
6   Furnace 2 [kW]                      503910 non-null  float64
7   Home office [kW]                    503910 non-null  float64
8   Fridge [kW]                         503910 non-null  float64
9   Wine cellar [kW]                   503910 non-null  float64
10  Garage door [kW]                    503910 non-null  float64
11  Kitchen 12 [kW]                     503910 non-null  float64
12  Kitchen 14 [kW]                     503910 non-null  float64
13  Kitchen 38 [kW]                     503910 non-null  float64
14  Barn [kW]                           503910 non-null  float64
15  Well [kW]                           503910 non-null  float64
16  Microwave [kW]                      503910 non-null  float64
17  Living room [kW]                    503910 non-null  float64
18  Solar [kW]                          503910 non-null  float64
19  temperature                          503910 non-null  float64
20  icon                                503910 non-null  object
21  humidity                            503910 non-null  float64
22  visibility                          503910 non-null  float64
23  summary                             503910 non-null  object
24  apparentTemperature                 503910 non-null  float64
25  pressure                           503910 non-null  float64
26  windSpeed                          503910 non-null  float64
27  cloudCover                         503910 non-null  object
28  windBearing                        503910 non-null  float64
29  precipIntensity                     503910 non-null  float64
30  dewPoint                           503910 non-null  float64
31  precipProbability                   503910 non-null  float64
dtypes: float64(28), object(4)
memory usage: 123.0+ MB
```

```
In [8]: #drop row with Nan in stipulated column
df = df.dropna(subset=["House overall [kW]", "temperature", "humidity",
                      "apparentTemperature",
                      "pressure", "Fridge [kW]", "Microwave [kW]",
                      "Dishwasher [kW]",
                      "Garage door [kW]",
                      "Living room [kW]"])

df
```

Out[8]:

	time	use [kW]	gen [kW]	House overall [kW]	Dishwasher [kW]	Furnace 1 [kW]	Furnace 2 [kW]	Home office [kW]	
0	1451624400	0.932833	0.003483	0.932833	0.000033	0.020700	0.061917	0.442633	0.
1	1451624401	0.934333	0.003467	0.934333	0.000000	0.020717	0.063817	0.444067	0.
2	1451624402	0.931817	0.003467	0.931817	0.000017	0.020700	0.062317	0.446067	0.
3	1451624403	1.022050	0.003483	1.022050	0.000017	0.106900	0.068517	0.446583	0.
4	1451624404	1.139400	0.003467	1.139400	0.000133	0.236933	0.063983	0.446533	0.
...
503905	1452128305	1.601233	0.003183	1.601233	0.000050	0.085267	0.642417	0.041783	0.
503906	1452128306	1.599333	0.003233	1.599333	0.000050	0.104017	0.625033	0.041750	0.
503907	1452128307	1.924267	0.003217	1.924267	0.000033	0.422383	0.637733	0.042033	0.
503908	1452128308	1.978200	0.003217	1.978200	0.000050	0.495667	0.620367	0.042100	0.
503909	1452128309	1.990950	0.003233	1.990950	0.000050	0.494700	0.634133	0.042100	0.

503910 rows × 32 columns



```
In [9]: #Split data into features for training
X = df[["temperature", "humidity", "apparentTemperature", "pressure",
        "Fridge [kW]", "Microwave [kW]", "Dishwasher [kW]",
        "Garage door [kW]", "Living room [kW]"]]

Y = df[["House overall [kW]"]]
```

In [10]:

```

#Train and Test splitting
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size = 0.2,
                                                    random_state=42)

#Drop Nan rows values
X_train = X_train.dropna()
Y_train = Y_train.loc[X_train.index]

X_test = X_test.dropna()
Y_test = Y_test.loc[X_test.index]

#scale the features using standardscaler
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

#MLP
MLP = MLPRegressor(hidden_layer_sizes=(64, 32), activation='relu', max_iter=2000,
                    learning_rate_init=0.001, random_state=42)

#Training of Model
MLP.fit(X_train, Y_train)

```

C:\Users\user\anaconda3\Lib\site-packages\sklearn\neural_network_multilayer_perceptron.py:1625: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
y = column_or_1d(y, warn=True)
```

Out[10]:

MLPRegressor	
MLPRegressor(hidden_layer_sizes=(64, 32), max_iter=2000, random_state=42)	

In [11]:

```

#first prediction
X_prediction = pd.DataFrame({
    "temperature" : [36.17],
    "humidity" : [0.20],
    "apparentTemperature" : [29.22],
    "pressure" : [1016.90],
    "Fridge [kW]" : [0.123540],
    "Microwave [kW]" : [0.004065],
    "Dishwasher [kW]" : [0.000020],
    "Garage door [kW]" : [0.013081],
    "Living room [kW]" : [0.00162]
})

predict = MLP.predict(X_prediction)
predict

```

C:\Users\user\anaconda3\Lib\site-packages\sklearn\base.py:457: UserWarning: X has feature names, but MLPRegressor was fitted without feature names

```
warnings.warn(
```

Out[11]: array([373.249988])

```
In [12]: # make predictions using the MLP regressor
```

```
Y_pred = MLP.predict(X_test)
```

```
#calculate Mean absolute error (MAE)
```

```
MAE = mean_absolute_error(Y_test,Y_pred)
```

```
print("Mean Absolute Error (MLP):",MAE)
```

Mean Absolute Error (MLP): 0.2845635740535989

```
In [13]: # visualize the results using a scatter plot
```

```
ya = Y_test
```

```
yp = MLP.predict(X_test)
```


```
plt.scatter(ya,yp)
```

```
plt.xlabel("Actual Values MLP")
```

```
plt.ylabel("Predicted Values MLP")
```

```
plt.title("Actual vs. Predicted values")
```

```
plt.show()
```




```
In [14]: #training of random forst model
from sklearn.ensemble import RandomForestRegressor
regressor = RandomForestRegressor(n_estimators = 10, random_state=0)
regressor.fit(X_train, Y_train)
```

C:\Users\user\anaconda3\Lib\site-packages\sklearn\base.py:1151: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
return fit_method(estimator, *args, **kwargs)

```
Out[14]: RandomForestRegressor
RandomForestRegressor(n_estimators=10, random_state=0)
```


```
In [15]: # make predictions using the Random forrest regressor
Y_pred_RF = regressor.predict(X_test)

#calculate Mean absolute error (MAE)
MAE = mean_absolute_error(Y_test,Y_pred_RF)
print("Mean Absolute Error (RF):",MAE)
```

Mean Absolute Error (RF): 0.1366152315021199

```
In [16]: # visualize the results using a scatter plot
ya_RF = Y_test
yp_RF = regressor.predict(X_test)

plt.scatter(ya_RF,yp_RF)
plt.xlabel("Actual Values of Random forest model")
plt.ylabel("Predicted Values of Random forest model ")
plt.title("Actual vs. Predicted values")
plt.show()
```



In []: